

1. Comparing Containers

Containers type

1. cloud
2. HPC
3. IoT

Containers alternative

4. unikernel

Cloud

Containers power microservices architecture that have their benefits for application services. Any application built with microservices-based methodology tends to have a chunk of small services that can be packaged in containers. All those services maintain their life cycle along with service-specific requirements of independent development, granular scaling and patching and fault remediation.

Containers for Cloud

- **consistent environment**
- **run anywhere**
- **isolation**
- **low overhead**
- **small startup time**
- **service based**
- ...

HPC

HPC leverages distributed compute and storage resources to solve complex issues with large data volumes. HPC clusters are commonly known as Supercomputers. Complex algorithms are used on large data sets to generate insights. HPC systems use a large set of CPU or GPU in parallel architecture that creates enough of a computing resource pool to execute complex mathematical algorithms.

HPC are mainly used for scientific research and big data analysis.

Containers for HPC

- **modularity:** HPC applications requires software dependencies that are numerous, complex, unusual, differently configured, or simply newer/older than what the center provides
- **scaling:** react to situations when there is a spike in the data processing, scaling the number of container to tackle such spikes

- **portability:** run the application on workstations and other test and development system not managed by the center
- **consistency:** environments can be easily, reliably, and verifiably reproduced in the future
- **security:** exploit Linux user namespaces to run containers with no privileged operations in a multi-tenant system
- **native performance:** exploit dedicated hardware with less possible overhead
- **support for GPU**
- **support for parallel filesystems and diskless computing nodes**
- **support for workload manager and job scheduler**

Docker and HPC issues

- Docker emulates a virtual machine in many aspects, users can escalate to root
- Docker uses a root owned daemon that users can control by means of a writable socket
- can not limit access to local file systems
- no native support for GPU
- no native support for MPI
- incompatibilities with existing scheduling and resource manager paradigms
- HPC is not the appropriate use-case or interest for the Docker community

Parallelism with HPC containers

Jobs are dispatched to containers using a workload manager / resource manager (such as Slurm, SGE). Some container technologies provide native support for MPI.

IoT

IoT, refers to network of physical devices collecting and sharing data.

Containers for IoT

- **limited network access:** use less bandwidth, downloading only the necessary part
- **different devices environments:** support for wide variety of chipset architectures. It also simplify the software configuration
- **minimal hardware resources:** use the resources (memory and computing power) of the device more conservatively. Configuration processing is minimal.
- **use smaller binaries:** reduce the resource consumption

Docker and IoT issues

- not suitable for all IoT software deployment
- not compatible with all of the hardware used by IoT
- heavyweight container runtime not suitable for IoT devices with few hardware resources
- not designed to reduce the network bandwidth consumption

Unikernel

A unikernel is a type of microservices environment that contains absolutely everything required to run a particular piece of software—including not just the software code itself, but also the operating system code necessary to host it. Plus, everything that is not strictly necessary for hosting the app is stripped out of the unikernel.

Unikernel Advantages

- self-hosted, portable and minimalist
- very small overhead
- suitable for cloud application: unikernel as an alternative to VMs
- suitable for IoT: unikernel provides everything needed to deploy the software for an IoT device
- suitable for device drivers: drivers can be supplied on an as-needed basis inside portable environments
- suitable for on-demand computing: fast to boot

References

- Containers and HPC: Mutually Beneficial
- Intel HPC developer conference
- <https://sylabs.io/>
- Singularity on HPC
- <https://hpc.github.io/charliecloud/index.html>
- Why the Internet of Things Needs Docker Containers
- Unikernels Use Cases: IoT, the Cloud and More

2.1 Comparing Solutions - HPC

Comparison

	Open Source	First Release	Last Release	Stars	Test	Issues (open/closed)
Singularity	Yes	Oct. 2015	Feb. 2020	1.6k	Ok	370 / 2250
Shifter	Yes	Mar. 2015	Jan. 2020	280	Not Ok	33 / 61
CharlieCloud	Yes	Jun. 2015	Mar. 2020	171	Ok	117 / 306
Sarus	Yes	Nov. 2018	Feb. 2020	32	Ok	5 / 1

2.1.1 Singularity

- <https://sylabs.io/>
- **Report**

Introduction

Singularity is an open source container platform designed to be simple, fast, and secure. Singularity is optimized for compute focused enterprise and HPC workloads, allowing untrusted users to run untrusted containers in a trusted way.

- Secure, single-file based container format
- Support for data-intensive workloads
- Extreme mobility
- Compatibility
- Simplicity
- Security
- User groups

Reference

- <https://github.com/sylabs/singularity>

Requirements and Permissions

- Singularity requires ~140MiB disk space once compiled and installed.
- 2GB of RAM is recommended when building from source Full functionality of Singularity requires that the kernel supports:
- overlayFS mounts (minimum kernel ≥ 3.18): required for full flexibility in bind mounts to containers, and to support persistent overlays for writable containers
- unprivileged user namespaces (minimum kernel ≥ 3.8 , ≥ 3.18 recommended): required to run containers without root or setuid privilege

Reference

- system requirements

Standards

- OCI compliant from 3.1.0
- OpenPGP standard to create and manage SIF signatures

Reference

- Singularity 3.1.0 brings in Full OCI Compliance

Images

Based on SIF: single file based, compressed, cryptographically signed, trusted, and immutable

- can be pulled from:
 - Singularity Library
 - Docker Hub
 - Singularity Hub
 - local
 - directory
 - definition file

Some comparisons between images from DockerHub and Singularity Library:

Image	Version	Arch	Singularity	DockerHub
Alpine	3.7	amd64	1.99 MB	2.01 MB
Ubuntu	18.04	amd64	35.38 MB	25.49 MB
CentOS	7	amd64	79.08 MB	72.27 MB

Reference

- SIF
- Build a Container

Performance

// TODO

Security

Singularity uses a number of strategies to provide safety and ease-of-use on both single-user and shared systems. 1. the user inside a container is the same as the user who ran the container. This means access to files and devices from the container is easily controlled with standard POSIX permissions

The Singularity Runtime enforces a unique security model that makes it appropriate for untrusted users to run untrusted containers safely on multi-tenant resources. When you run a container, the processes in the container will run as your user account. Singularity dynamically writes UID and GID information to the appropriate files within the container, and the user remains the same inside and outside the container, i.e., if you're an unprivileged user while entering the container you'll remain an unprivileged user inside the container.

2. container filesystems are mounted nosuid and container applications run with the PR_NO_NEW_PRIVS flag set.

Additional blocks are in place to prevent users from escalating privileges once they are inside of a container. The container file system is mounted using the nosuid option, and processes are started with the PR_NO_NEW_PRIVS flag set. This means that even if you run

sudo inside your container, you won't be able to change to another user, or gain root privileges by other means. This approach provides a secure way for users to run containers and greatly simplifies things like reading and writing data to the host system with appropriate ownership.

3. the Singularity Image Format (SIF) supports encryption of containers, as well as cryptographic signing and verification of their content

You generally do not need admin/sudo to use Singularity containers but you do however need admin/root access to install Singularity and for some container build functions (for example, building from a recipe, or a writable image). This then defines the work-flow to some extent. If you have a container (whether Singularity or Docker) ready to go, you can run/shell/import without root access. If you want to build a new Singularity container image from scratch it must be built and configured on a host where you have root access (this can be a physical system or on a VM). And of course once the container image has been configured it can be used on a system where you do not have root access as long as Singularity has been installed there.

4. SIF containers are immutable and their payload is run directly, without extraction to disk. This means that the container can always be verified, even at runtime, and encrypted content is not exposed on disk. Restrictions can be configured to limit the ownership, location, and cryptographic signatures of containers that are permitted to be run

A SIF file is an immutable container image that packages the container environment into a single file. SIF supports security and integrity through the ability to cryptographically sign a container, creating a signature block within the SIF file which can guarantee immutability and provide accountability as to who signed it. Singularity follows the OpenPGP standard to create and manage these signatures, and the keys used to create them

Reference

- Security in Singularity
- Singularity Security

Available Tools

- Support for Docker and OCI
- MPI
- GPU support
- plugins
- cloud library
- Singularity CRI: consists of two separate services: runtime and image, each of which implements K8s RuntimeService and ImageService respectively

Limits

// TODO

2.1.2 Shifter

- <https://github.com/NERSC/shifter>
- **Report**

Introduction

Shifter enables container images for HPC. Shifter allows an HPC system to efficiently and safely allow end-users to run a docker image. Shifter consists of a few moving parts

1. a utility that typically runs on the compute node that creates the run time environment for the application
2. an image gateway service that pulls images from a registry and repacks it in a format suitable for the HPC system (typically squashfs)
3. example scripts/plugins to integrate Shifter with various batch scheduler systems

These components are:

1. Shifter Runtime: Instantiates images securely on compute resources
2. Shifter Image Gateway: Imports and converts images from DockerHub and Private Registries
3. Work Load Manager Integration: Integrates Shifter with WLM

Design Goals:

- User independence: require no administrator assistance to launch an application inside an image
- Shared resource availability (e.g., file systems and network interfaces)
- Leverages or integrates with public image repos (i.e. DockerHub)
- Seamless user experience
- Robust and secure implementation

Reference

- [Index](#)

Requirements and Permissions

- requires the installation of the Image Manager and the Shifter Runtime. The image manager doesn't really do any real work on its own, and the image worker uses only user-space tools to construct images (in the default configuration) so they do not need to run with root privilege

Reference

- [Installation](#)




Figure 1: Shifter Architecture

Standards

- Docker-compatible container platform

Images

- Docker Images
- uses squashfs

Shifter works by converting Docker images to a common format that can then be efficiently distributed and launched on HPC systems. The user interface to shifter enables a user to select an image from their DockerHub account or the NERSC private registry and then submit jobs which run entirely within the container.

Shifter works by enabling users to pull images from a DockerHub or private docker registry. An image manager at NERSC then automatically converts the image to a flattened format that can be directly mounted on the compute nodes. This image is copied to the Lustre scratch filesystem (in a system area). The user can then submit jobs specifying which image to use. Private images are only accessible by the user that authenticated and pulled them, not by the larger community. In the job the user has the ability to either run a custom batch script to perform any given command supported by the image, or if a Docker entry-point is defined, can simply execute the entry-point

Reference

- Building Shifter Images

Performance

// TODO

Security

Shifter use a great deal of root privilege to setup the container environment. The “shifter” executable is setuid-root, and when run with batch integration the setupRoot/unsetupRoot utilities must be run as root. We are working to reduce the privilege profile of starting Shifter containers to reduce risk as much as possible.

Once a process is launched into the container, processes are stripped of all privilege, and should not be able to re-escalate afterwards.

Shifter enables User Defined Image environment containers. To do this while optimizing I/O on the compute nodes it does require performing several privileged operations on the execution node, both privilege to mount filesystems and rewrite the user space, and privilege to manipulate devices on the system.

Furthermore, because the environment is *user defined*, it is possible that a user could include software which could generate security vulnerabilities if a privileged process accesses such software or resources.

Reference

- Security

Available Tools

- native GPU support: automatic import of host’s CUDA driver and devices
- native MPI support:
 - transparently swap container’s MPI libraries with the host’s at runtime
 - enables full performance from vendor-specific implementations
 - relies on MPICH ABI compatibility
- Shifter is distributed with a SPANK plugin for SLURM.

Reference

- shifter docker containers for hpc

Limits

// TODO

2.1.3 Charliecloud

- <https://hpc.github.io/charliecloud/index.html>
- **Report**

Introduction

Charliecloud provides user-defined software stacks (UDSS) for high-performance computing (HPC) centers. This “bring your own software stack” functionality addresses needs such as:

- software dependencies that are numerous, complex, unusual, differently configured, or simply newer/older than what the center provides;
- build-time requirements unavailable within the center, such as relatively unfettered internet access;
- validated software stacks and configuration to meet the standards of a particular field of inquiry;
- portability of environments between resources, including workstations and other test and development system not managed by the center;
- consistent environments, even archivally so, that can be easily, reliably, and verifiably reproduced in the future; and/or
- usability and comprehensibility.

Reference

- [What is Charliecloud?](#)

Requirements and Permissions

- running the Docker daemon and executing Docker commands require privileged access, this happens on user-managed resources;
- nothing in the Charliecloud workflow requires privileged or trusted processes or daemons on center-managed resources. All privileged steps take place on user systems, and the scripts escalate with sudo as needed

Reference

- [paper](#)

Standards

- OCI compliant is in beta with ch-run-oci

Reference

- [ch-run-oci](#)

Images

- Container images can be built using Docker or anything else that can generate a standard Linux filesystem tree

Performance

- Charliecloud containers impose minimal performance penalty, because the guest is using the same kernel and devices as the host
- Performance is the same as native in tests because minimal isolation yields direct access to all resources

Reference

- [paper](#)

- paper

Security

Charliecloud uses Linux user namespaces to run containers with no privileged operations or daemons and minimal configuration changes on center resources. This simple approach avoids most security risks while maintaining access to the performance and functionality already on offer.

Charliecloud relies on two things to maintain security:

1. the Linux kernel to enforce access control and other aspects of security
2. the extension that user namespaces will ensure that guest UIDs are an illusion

Reference

- Interacting with the host
- paper

Available Tools

- native MPI support
- inject GPU libraries to a container with ch-fromhost

Limits

// TODO

2.1.4 Sarus

- <https://github.com/eth-cscs/sarus>
- **Report**

Introduction

Sarus is a software to run Linux containers on High Performance Computing environments. Its development has been driven by the specific requirements of HPC systems, while leveraging open standards and technologies to encourage vendor and community involvement.

Key features:

- Spawning of isolated software environments (containers), built by users to fit the deployment of a specific application
- Security oriented to HPC systems
- Extensible runtime by means of OCI hooks to allow current and future support of custom hardware while achieving native performance
- Creation of container filesystems tailored for diskless nodes and parallel filesystems

- Compatibility with the presence of a workload manager
- Compatibility with the Open Container Initiative (OCI) standards:
 - Can pull images from registries adopting the OCI Distribution Specification or the Docker Registry HTTP API V2 protocol
 - Can import and convert images adopting the OCI Image Format
 - Sets up a container bundle complying to the OCI Runtime Specification
 - Uses an OCI-compliant runtime to spawn the container process

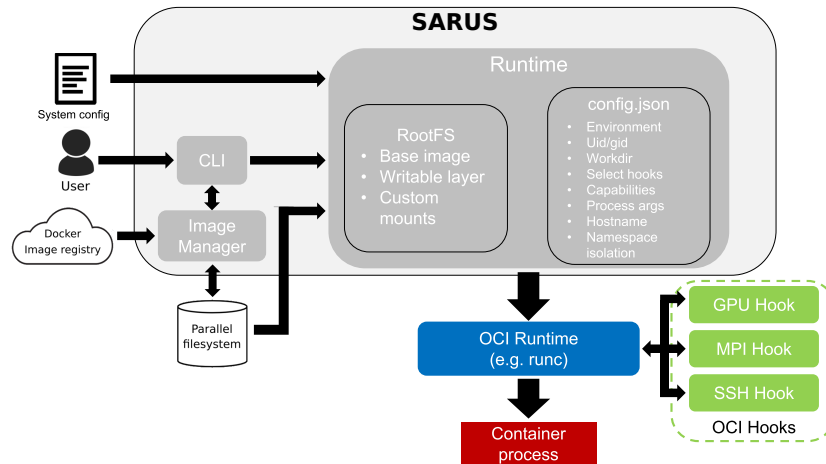


Figure 2: Sarus architecture

Reference

- Sarus - An OCI-compatible container engine for HPC

Requirements and Permissions

- Sarus is installed on the HPC and it requires root in order to set Sarus as a root-owned SUID program

Standards

- designed around the specifications of the OCI

Images

- Docker Images
- uses squashfs

Performance

- the presented results show no performance degradation when comparing the natively compiled versions with their containerized counterparts

Reference

- Sarus: Highly Scalable Docker Containers for HPC Systems

Security

- Sarus must run as a root-owned SUID executable and be able to achieve full root privileges to perform mounts and create namespaces
- Write/read permissions to the Sarus's centralized repository. The system administrator can configure the repository's location through the centralizedRepositoryDir entry in sarus.json
- Write/read permissions to the users' local image repositories. The system administrator can configure the repositories location through the localRepositoryBaseDir entry in sarus.json
- The configuration script needs to run with root privileges in order to set Sarus as a root-owned SUID program

Reference

- Review permissions required by Sarus

Available Tools

OCI Hooks:

- Native MPICH-Based MPI Support
- NVIDIA GPU Support
- Slurm Scheduler Synchronization

Reference

- Sarus: Highly Scalable Docker Containers for HPC Systems

Limits

// TODO

Reference

- State of HPC containers